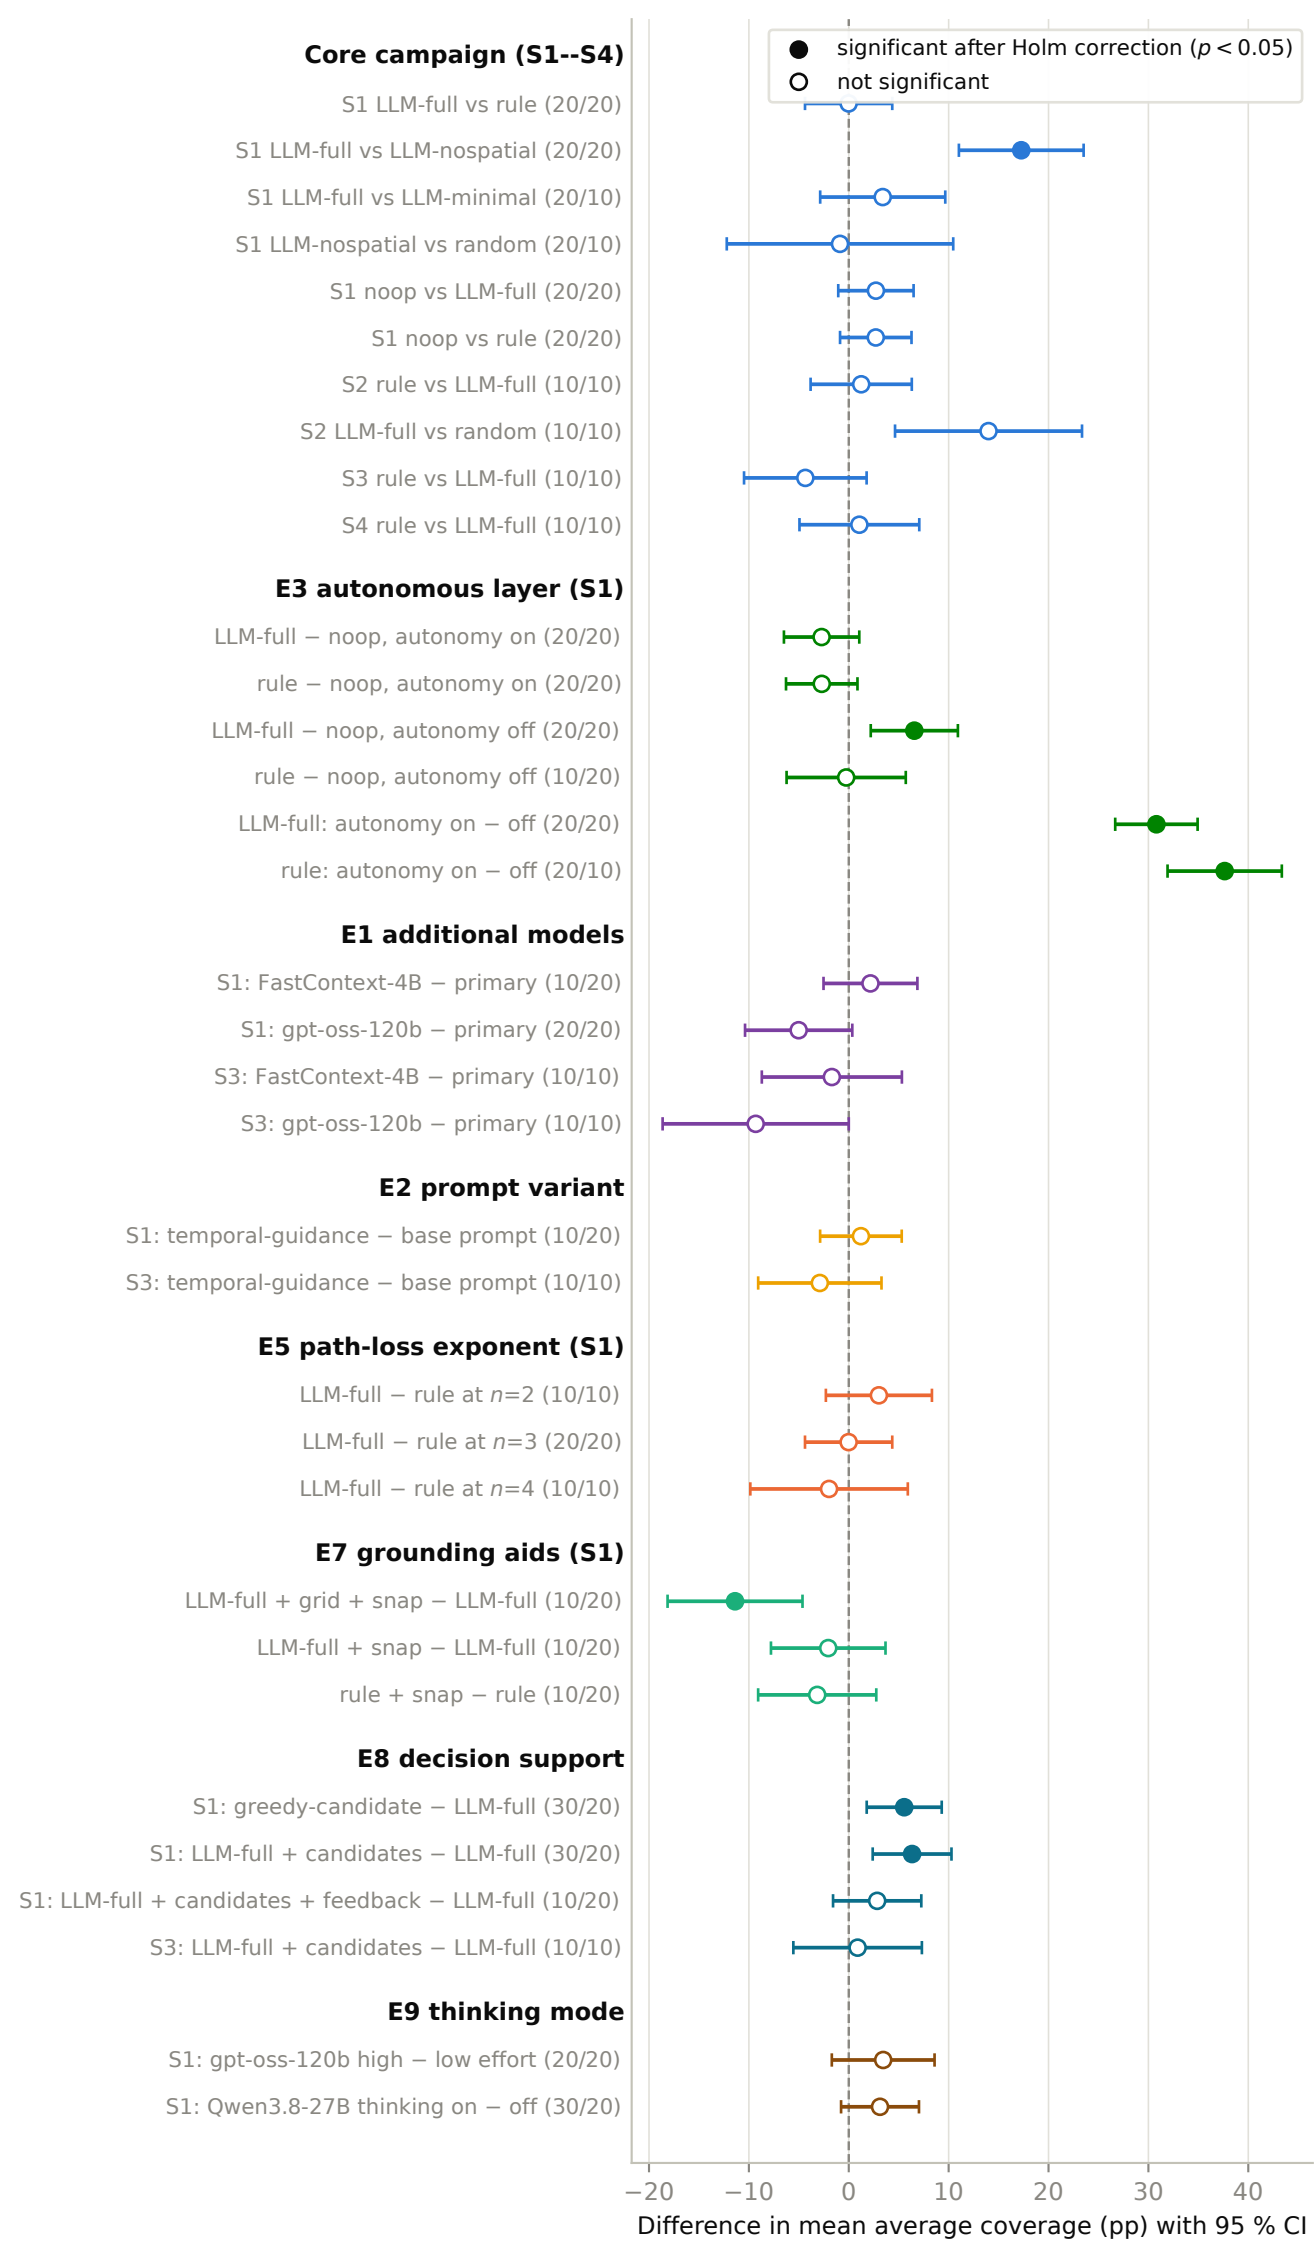
significant after Holm correction ($p < 0.05$)
not significant
Core campaign (S1--S4)
S1 LLM-full vs rule (20/20)
S1 LLM-full vs LLM-nospatial (20/20)
S1 LLM-full vs LLM-minimal (20/10)
S1 LLM-nospatial vs random (20/10)
S1 noop vs LLM-full (20/20)
S1 noop vs rule (20/20)
S2 rule vs LLM-full (10/10)
S2 LLM-full vs random (10/10)
S3 rule vs LLM-full (10/10)
S4 rule vs LLM-full (10/10)
E3 autonomous layer (S1)
LLM-full − noop, autonomy on (20/20)
rule − noop, autonomy on (20/20)
LLM-full − noop, autonomy off (20/20)
rule − noop, autonomy off (10/20)
LLM-full: autonomy on − off (20/20)
rule: autonomy on − off (20/10)
E1 additional models
S1: FastContext-4B − primary (10/20)
S1: gpt-oss-120b − primary (20/20)
S3: FastContext-4B − primary (10/10)
S3: gpt-oss-120b − primary (10/10)
E2 prompt variant
S1: temporal-guidance − base prompt (10/20)
S3: temporal-guidance − base prompt (10/10)
E5 path-loss exponent (S1)
LLM-full − rule at n=2 (10/10)
LLM-full − rule at n=3 (20/20)
LLM-full − rule at n=4 (10/10)
E7 grounding aids (S1)
LLM-full + grid + snap − LLM-full (10/20)
LLM-full + snap − LLM-full (10/20)
rule + snap − rule (10/20)
E8 decision support
S1: greedy-candidate − LLM-full (30/20)
S1: LLM-full + candidates − LLM-full (30/20)
S1: LLM-full + candidates + feedback − LLM-full (10/20)
S3: LLM-full + candidates − LLM-full (10/10)
E9 thinking mode
S1: gpt-oss-120b high − low effort (20/20)
S1: Qwen3.8-27B thinking on − off (30/20)
−20
−10
0
10
20
30
40
Difference in mean average coverage (pp) with 95 % CI